Project Report:
Multi-Label Text Classification on the Reuters Corpus

Linyao Du, Aojie Yin
Deep Learning for NLP
University of Helsinki

1 Objective and Data Challenge

The objective of this project is to build a multi-label text classification system for the Reuters Corpus (RCV1).

The core challenge is extreme class imbalance. The dataset follows a severe long-tail distribution. A few dominant labels, such as `E21` (GOVERNMENT FINANCE/GVOTE-ELECTIONS), appear very frequently, while many rare categories have extremely few samples. This makes it easy for models to systematically ignore minority labels.

2 Data Pipeline

XML Parsing: Built a custom `BeautifulSoup` parser to extract the `headline` and `text` from 299,000 documents, processing everything in under an hour.

Label Binarization: Extracted topic codes and converted them into 0/1 vectors using `MultiLabelBinarizer`.

Sampling: Evaluated on a 50,000-record subset to allow for faster iteration without losing the overall feature distribution.

3 Experimental Setup

3.1 Baseline Model

A traditional machine learning pipeline was built as a control group:

Features: TF-IDF (top 20,000 words).

Model: Logistic Regression with `OneVsRest` strategy.

Result: Achieved an F1 score of **0.8064**.

3.2 Deep Learning Approach

The proposed neural network is **DistilBERT**, selected because it is 40% smaller than full BERT but retains 97% of its performance, making it highly efficient. The model used `BCEWithLogitsLoss`, which is the industry standard for independent multi-label tasks.

Three strategies were implemented to optimize performance:

1. **Threshold Sweeping:** The default 0.5 threshold was too conservative, yielding high Precision (91.9%) but low Recall (84.1%).
2. **Stage-2 Fine-Tuning (Annealing):** To break convergence plateaus, the learning rate was dropped from 3e-5 to 1e-5 for an extra epoch to carefully polish the weights.
3. **Context Expansion:** Increased the `max_length` from 256 to 512 tokens to capture terminal clues at the end of long news documents. Batch size was reduced to 8 to prevent GPU Out-of-Memory (OOM) errors.

4 Results

4.1 Threshold Analysis

Before further optimizing the model, we evaluated the precision-recall tradeoff at different decision thresholds for the initial DistilBERT model (Model 1). As shown in Table 1, the default threshold of 0.5 yielded high precision but suppressed recall. Shifting the threshold to 0.40 caused a significant Recall surge (to 86.44%), maximizing the overall F1 score. This 0.40 threshold was then adopted for all subsequent models.

Threshold	Micro-F1	Precision	Recall
0.30	0.8797	0.8718	0.8877
0.35	0.8815	0.8868	0.8762
0.40 (best)	**0.8816**	**0.8993**	**0.8645**
0.45	0.8805	0.9099	0.8530
0.50 (default)	0.8785	0.9189	0.8415
0.55	0.8758	0.9272	0.8299

Table 1: Precision-recall tradeoff across decision thresholds.

4.2 Ablation Study

Table 2 tracks the progressive improvements on the validation set as each optimization (threshold sweeping, annealing, and context expansion) was added.

Model	Configuration	Thr.	F1
Baseline	TF-IDF + LR	0.6	0.8064
Model 1	DistilBERT (len=256)	0.5	0.8785
Model 1	Threshold sweeping	0.4	0.8815
Model 2	Stage-2 (lr=1e-5)	0.4	0.8847
Model 3	**DistilBERT (len=512)**	**0.4**	**0.8902**

Table 2: Ablation study results on the validation set.

4.3 Final Blind Test Results

The final model (Model 3) was evaluated against the official ground truth on the held-out test set.

Metric	Validation	Blind Test	Gap
Micro-F1	0.8902	0.8819	-0.0083
Precision	0.8975	0.8747	-0.0228
Recall	0.8831	0.8892	+0.0061

Table 3: Validation vs. blind test performance.

Performance Stability: The generalization gap between the validation F1 (0.8902) and blind test F1 (0.8819) is less than 1%. This demonstrates zero overfitting despite the increased model complexity of the 512-token length.

5 Error Analysis: Semantic Overlap & Model Bias

A manual inspection of misclassified cases (e.g., test line 1667 regarding "Bre-X Minerals") revealed a systematic bias. The first 5 errors examined all showed the exact same pattern.

Example Excerpt (Line 1667):

> *"...A flurry of news releases, shareholder lawsuits and a deepening wrangle between Bre-X and TSE regulators... Despite strong wording contained in four class action lawsuits filed in the United States..."*

Ground Truth: C11 (Strategy/Plans)

Prediction: C11 (Strategy) + **C12 (Legal/Judicial)**

Why the Model Fails: The article discusses a corporate strategy in response to a crisis, but contains the background legal terms highlighted in the excerpt above. The model got distracted by these local keywords.

This failure stems from three root causes:

1. **Feature Over-indexing:** The model gives too much weight to heavy legal terms and treats them as primary topics instead of secondary context.
2. **Spurious Correlation:** The model learned a false rule during fine-tuning (e.g., "if a lawsuit is mentioned, it must be C12").
3. **Global Intent Deficit:** Despite having a 512-token context, DistilBERT lacks the complex reasoning needed to separate the core strategic narrative from the background legal environment. Keyword frequency is not the same as narrative intent.

6 Lessons Learned & Possible Improvements

6.1 Lessons Learned

Context is King: Breaking the 256-token limit to 512 was the most significant boost, proving that news summaries often hide critical cues in closing paragraphs.

Threshold Engineering: The default 0.5 boundary is rarely optimal for imbalanced data. Dynamic sweeping is essential.

Annealing Efficiency: Dropping the learning rate in Stage-2 is a resource-efficient way to fine-polish weights.

Engineering Integrity: A strong, optimized baseline is the only way to prove the true Return on Investment (ROI) of Deep Learning.

6.2 Possible Future Improvements

Architecture & Loss:

Implement **Focal Loss** to penalize errors on "easy" samples and force the model to learn "hard" rare labels.

Upgrade to **RoBERTa-Large** for better logical reasoning in ambiguous cases (like Strategy vs. Legal).

Data & Scaling:

Unlock the full 300k corpus to increase the exposure of rare labels during training.

Explore **Zero-shot prompting** with Large Language Models (LLMs) to handle extreme long-tail categories that have zero training samples.